

Deploying and Evaluating LLMs to Program Service Mobile Robots

Francesca Lucchetti¹, Zichao Hu², Anders Freeman³, Sadanand Modak², Yash Saxena², Luisa Mao²,
Claire Schlesinger¹, Arjun Guha¹, Joydeep Biswas²

I. INTRODUCTION

Recent advancements in large language models (LLMs) have spurred interest in using them for generating robot programs from natural language, with promising initial results [1], [2], [3]. We investigate the use of LLMs to generate programs for *service mobile robots* leveraging mobility, perception, and human interaction skills, where accurate *sequencing and ordering of actions* is crucial for success. We contribute CodeBotler, an open-source tool to program service mobile robots from natural language, and RoboEval, a benchmark to evaluate the correctness and robustness of generated programs. We release our code and benchmark at <https://amrl.cs.utexas.edu/codebotler/>.

II. CODEBOTLER

CodeBotler is an open-source tool to generate general-purpose service robot programs from natural language, and to deploy such programs on general-purpose autonomous mobile robots. It performs program generation with an *embedded domain-specific language* (eDSL) via few-shot prompting in a programming language that LLMs are already adept in: Python. By abstracting robot skills in the eDSL, we release CODEBOTLER, a robot-agnostic deployment system for executing generated programs on any general-purpose mobile robot. By embedding the eDSL in Python, CODEBOTLER drastically reduces the number of syntax and run-time errors of generated code.

III. ROBOEVAL

Existing code completion benchmarks tackle simple data processing functions [4] or low-level robot skills [3], which are amenable to simple input-output unit tests. However, code generation for general-purpose service robot programs cannot be evaluated just on input/output sequences. For example, given a task “*Check how many conference rooms have no markers*”, it is insufficient to just test whether the LLM-generated program executes to state the correct answer — to ensure correctness, a correct program must first visit all conference rooms and check for markers there before arriving at the result. We also observe that there are significant variations in the correctness of generated programs with small variations in the phrasing of the natural language task descriptions [5]. We thus introduce three key ideas to test for both correctness and robustness of LLM-generated robot programs: 1) we evaluate the *execution traces* of programs; 2) we check whether the execution traces satisfy *temporal logic* properties that encode correctness for each task; and 3)

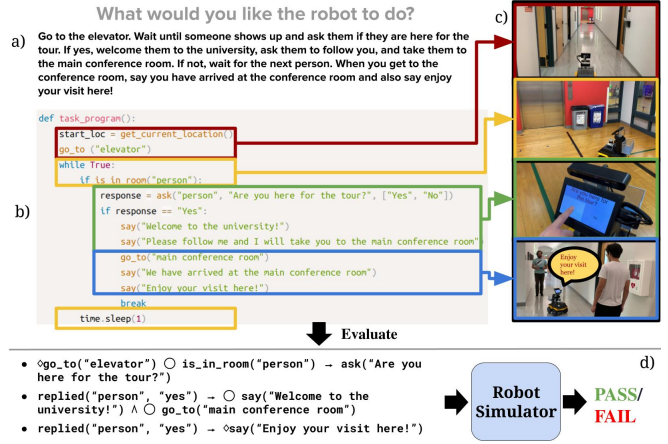


Fig. 1: CODEBOTLER converts a) natural language task prompts to b) python programs and c) executes on a real robot. d) ROBOEVAL verifies results via LTL checking of simulation traces.

we *vary the prompts* and to test for robustness. Combining these ideas, we release a preview of ROBOEVAL with five tasks, each with multiple world states, multiple prompts, and temporal logic formulas to check correctness. Fig. 2 shows the results of three LLMs (PaLM, GPT-3.5, and StarCoder) on the ROBOEVAL benchmark.




Fig. 2: ROBOEVAL pass@1 results for PaLM (TEXT-BISON-001), StarCoder, and GPT3.5 (TEXT-DAVINCI-003).

REFERENCES

- [1] W. Huang, P. Abbeel, D. Pathak, and I. Mordatch, “Language models as zero-shot planners: Extracting actionable knowledge for embodied agents,” in *ICML 2022*.
- [2] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, “ProgPrompt: Generating Situated Robot Task Plans using Large Language Models,” in *ICRA 2023*.
- [3] J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, “Code as policies: Language model programs for embodied control,” in *arXiv:2209.07753*, 2022.
- [4] M. Chen, J. Tworek, H. Jun, Q. Yuan, *et al.*, “Evaluating large language models trained on code,” *arXiv:2107.03374*, 2021.
- [5] H. M. Babe, S. Nguyen, Y. Zi, A. Guha, M. Q. Feldman, and C. J. Anderson, “StudentEval: A benchmark of student-written prompts for large language models of code,” *arXiv:2306.04556*, 2023.

¹Northeastern Univ., ²Univ. of Texas at Austin, ³Wellesley College